

# Dijkstra's Algorithm (MPI & OpenMP)

## Introduction

The purpose of this assignment is to improve your familiarity with MPI and OpenMP by implementing Dijkstra's algorithm.

## Dijkstra's algorithm

You need to write two programs, *dijkstra\_mpi* and *dijkstra\_omp*, that will take as input a single file (an adjacency matrix) and a source vertex ID and output to a file the weight of a shortest path from the source vertex to every other vertex. As the names suggest, the first program will use MPI and the second OpenMP for parallelizing the algorithm. The input file will consist of  $n+1$  lines, with the first line containing only the number  $n$ , and each other line containing  $n$  values (i.e., a dense matrix), with each value representing the weight of the edge. You should assume that the values in the input file are represented as quad-precision floats. An edge that does not exist in the graph is represented by the value 0. Your output file should also follow this convention.

For example, an adjacency matrix for an undirected graph with five vertices and nine edges would be stored in a file as

```
5
inf  0.2  inf  3.7  0.4
2.1  inf  3.9  1    inf
2    0.01 inf  0.7  6
9    1.2  2.5  inf  inf
3    3    7    4    inf
```

## Parallelization strategy

Feel free to try any parallelization strategy you wish, including strategies we discussed in class. The only requirement is that the input graph should be split and distributed to processes. The code should work for very large graphs that do not fit in the memory of a single node as long as they fit in the aggregate memory of the set of nodes tasked to execute the program.

## Output

The output file consists of the shortest path weight vector written to a text file, one value per line. In other words, the  $k$ th line in the file has the weight of the shortest path from the source to the  $k$ th node in the graph. The weight from the source to itself should be marked as `inf`.

## Baseline code

Serial baseline code in C is provided for. A Makefile is provided with the code. Execute 'make dijkstra' to compile the code, and 'make clean' to remove the executable and artifacts. Note that the Makefile already has code necessary to compile your *dijkstra\_mpi* and *dijkstra\_omp* files, assuming your source

files are called `dijkstra_mpi.c` and `dijkstra_omp.c`. If not, edit the appropriate variables in the Makefile with the name of your sources file(s).

## Testing

Test graphs can be generated by executing the Python script `make_data.py`. The only Python library that the script requires is Numpy and the script should run in both Python 2.x and Python 3.x environments. The largest file, *10000.graph*, contains 100 million numbers, representing a graph with 10,000 nodes. For each input graph, you can run the provided serial program to get a reference solution. You can then use the 'diff' Linux program to see if your solution is equivalent with the given solution for a given problem size.

Note that the evaluation may be done using a different set of input files than the ones in the test directory. As such, do not over-specify your code to work perfectly on those files and less so on others not in that set.

## What you need to turn in

1. The source code of your program.
2. A short report that includes the following:
  - a. A short description of how you went about parallelizing dijkstra's algorithm. One thing that you should be sure to include is how you decided what work each process would be responsible for doing.
  - b. Timing results for your parallel execution with the combination of 1, 2, 3 nodes and 1, 2, 4, 8, 16, and 28 processes per node (or threads in the case of the OpenMP program) for the search step (i.e., not including time for I/O). For each experiment, timing should be obtained as the average of the search times when searching for 20 randomly-selected sources. These results should be reported in a table as well as a graph or chart of some sort. The following is an explicit list of the problems to time for each of these process counts:  
  
100.graph  
500.graph  
1000.graph  
5000.graph  
10000.graph
  - c. A brief analysis of your results. Some things to consider might be:
    - i. How does the number of processes/threads and the number of nodes affect the runtime? Why do you think that is?
    - ii. How does the size of the problem affect the runtime? Why do you think that is?
    - iii. How well does your program scale, i.e., if you keep increasing the number of processes, do you think the performance will keep increasing at the same rate?

**Do NOT include the test input files** in your submission. They're over 0.5GB. You will lose points for including test files.

## Submission specifications

- A makefile must be provided to compile and generate the executable file.
- The executable files should be named 'dijkstra\_mpi' and 'dijkstra\_omp'.
- Your programs should take as arguments the input file name, source ID (0-based index), and the output file name. An optional number of threads parameter may also be specified. For example the program would be invoked as follows,

```
mpirun -n 28 ./dijkstra_mpi 10000.graph 0 test.out  
[or]
```

```
./dijkstra_omp 10000.graph 0 test.out 28
```

where 10000.graph is the input graph, 0 is the ID of the source node, test.out is the output file, and 28 is the number of threads.

- Your program MUST print to standard out ONLY the timing information and MUST use the same format as that provided in the baseline code function print\_time in dijkstra.c.
- All files (code + report) MUST be in a single directory and the directory's name MUST be your university student ID. Your submission directory MUST include at least the following files (other auxiliary files may also be included):

```
<Student ID>/dijkstra_mpi.c  
<Student ID>/dijkstra_omp.c  
<Student ID>/Makefile  
<Student ID>/report.pdf
```

- Submission MUST be a tar.gz archive.
- The following sequence of commands should work on your submission file:

```
tar xzvf <Student ID>.tar.gz  
cd <Student ID>  
make  
ls -ld dijkstra_mpi  
ls -ld dijkstra_omp
```

This ensures that your submission is packaged correctly, your directory is named correctly, your makefile works correctly, and your output executable files are named correctly. If any of these does not work, modify it so that you do not lose points. I can answer questions about correctly formatting your submission BEFORE the assignment is due. Do not expect questions to be answered the night it is due.

## Evaluation criteria

The goal for this assignment is for you to become more familiar with the OpenMP and MPI APIs and develop efficient parallel programs. As such, the following things will be evaluated:

1. follows the assignment directions,
2. solve the problem correctly,
3. do so in parallel,
4. achieve speedup:

- 5 / 10 points will be reserved for this criterion.
- points will be awarded by comparing the relative performance of your solution to the solution of the benchmark program provided.
- The speedups obtained will probably depend on the size of the input file. It is not expected that you get good speedups for small files, but you should be able to get good speedups for large files.